

First things first

ARIMA MODELS IN R



David Stoffer

Professor of Statistics at the University
of Pittsburgh


About Me

- Professor of Statistics




About Me

- Professor of Statistics
- Co-author of two texts on time series
- `astsa` package




Time Series Data - I

```
library(astsa)
plot(jj, main = "Johnson & Johnson Quarterly Earnings per Share", type = "c")
text(jj, labels = 1:4, col = 1:4)
```




Time Series Data - II

```
library(astsa)
plot(globtemp, main = "Global Temperature Deviations", type= "o")
```



Time Series Data - III

```
library(xts)
plot(sp500w, main = "S&P 500 Weekly Returns")
```



Time Series Regression Models

Regression: $Y_i = \beta X_i + \epsilon_i$, where ϵ_i is white noise

White Noise:

- independent normals with common variance
- is basic building block of time series

AutoRegression: $X_t = \phi X_{t-1} + \epsilon_t$ (ϵ_t is white noise)

Moving Average: $\epsilon_t = W_t + \theta W_{t-1}$ (W_t is white noise)

ARMA: $X_t = \phi X_{t-1} + W_t + \theta W_{t-1}$

Let's practice!
ARIMA MODELS IN R

Stationarity and nonstationarity

ARIMA MODELS IN R





David Stoffer

Professor of Statistics at the University
of Pittsburgh

Stationarity

A time series is stationary when it is "stable", meaning:

- the mean is constant over time (no trend)
- the correlation structure remains constant over time



Stationarity

Given data, x_1, \dots, x_n we can estimate by averaging

For example, if the mean is constant, we can estimate it by the sample average \bar{x}


Pairs can be used to estimate **correlation** on different lags:

$(x_1, x_2), (x_2, x_3), (x_3, x_4), \dots$ for lag 1

$(x_1, x_3), (x_2, x_4), (x_3, x_5), \dots$ for lag 2

Southern Oscillation Index


Reasonable to assume stationary, but perhaps some slight trend.



Southern Oscillation Index

To estimate autocorrelation, compute the correlation coefficient between the time series and itself at various lags.

Here you see how to get the correlation at lag 1 and lag 6.




Random Walk Trend

Not stationary, but differenced data are stationary.


$$X_t$$

globtemp



$$X_t - X_{t-1}$$

diff(globtemp)




Trend Stationarity

Stationarity around a trend, differencing still works!


$$X_t$$

chicken



$$X_t - X_{t-1}$$


diff(chicken)



Nonstationarity in trend and variability

First log, then difference


X_t



$\log(X_t)$



$\log(X_t) - \log(X_{t-1})$



Let's practice!

ARIMA MODELS IN R

Stationary time series: ARMA

ARIMA MODELS IN R



David Stoffer

Professor of Statistics at the University
of Pittsburgh

Wold Decomposition



Wold proved that any stationary time series may be represented as a linear combination of white noise:

$$X_t = W_t + a_1 W_{t-1} + a_2 W_{t-2} + \dots$$

For constants a_1, a_2, \dots

Any **ARMA** model has this form, which means they are suited to modeling time series.

Note: Special case of MA(q) is already of this form, where constants are 0 after q-th term.

Generating ARMA using `arima.sim()`

- Basic syntax:

```
arima.sim(model, n, ...)
```

- `model` is a list with order of the model as `c(p, d, q)` and the coefficients
- `n` is the length of the series

Generating and plotting MA(1)


Generate MA(1) given by

$$X_t = W_t + 0.9W_{t-1}$$

Generating and plotting MA(1)

Generate MA(1) given by

$$X_t = W_t + 0.9W_{t-1}$$



```
x <- arima.sim(list(order = c(0, 0, 1), ma = 0.9), n = 100)
plot(x)
```

Generating and plotting AR(2)


Generate AR(2) given by

$$X_t = -0.9X_{t-2} + W_t$$

Generating and plotting AR(2)

Generate AR(2) given by

$$X_t = -0.9X_{t-2} + W_t$$



```
x <- arima.sim(list(order = c(2, 0, 0), ar = c(0, -0.9)), n = 100)
plot(x)
```


Let's practice!
ARIMA MODELS IN R